

CSCI 4372/5372/6397: Data Clustering

Phase 3: Normalization and Initialization

Submission Deadline: **Wednesday, March 11 (23:59)**

Objective: Normalize your data and implement an alternative initialization method for k -means.

Part I – Normalization: In clustering applications, attribute normalization is a common preprocessing step necessary to prevent attributes with large ranges from dominating the distance computations and to avoid numerical instabilities. In this phase, you will implement the “min-max normalization” method to normalize your attributes to $[0,1]$ prior to clustering. You can easily find the description of this simple normalization method in any Data Mining book, or on the WWW.

Warning #1: Do **not** normalize your input data matrix across its rows; it should be normalized across its columns. In other words, each attribute of the data set should be normalized **independently** of the other attributes.

Bonus for CSCI 4372 Students [5 points]; Mandatory for CSCI 5372 and 6397 Students: In addition to “min-max normalization”, implement the “z-score normalization” method.

Hint #1: When implementing either normalization method, you should **avoid** divisions by zero.

Hint #2: You might want to dump the normalized data set to a file or the terminal to make sure that your normalization routine works properly.

Part II - Initialization: In Phase 2, you implemented the standard k -means algorithm. Unfortunately, k -means is highly sensitive to initialization. In other words, different initial centers can result in vastly different results. Numerous initialization methods have been proposed to address this problem. In this phase, you will implement another initialization method for k -means and compare it with the “random selection” method that you implemented in Phase 2. You will compare the two initialization methods with respect to the following factors:

- **Initial SSE:** This is the SSE value computed after the initialization phase, before the clustering phase. It gives us a measure of the effectiveness of an initialization method by itself.
- **Final SSE:** This is the SSE value computed after the clustering phase. It gives us a measure of the effectiveness of an initialization method when its output is refined by k -means. Note that this is the objective function of the k -means algorithm.
- **Number of Iterations:** This is the number of times k -means iterates until reaching convergence when initialized by a particular initialization method. It is an efficiency measure independent of programming language, implementation style, compiler, and CPU architecture.

In Phase 2, you implemented the “random selection” method. In this phase, you will implement the “random partition” method. In this method, starting from empty clusters, first assign each point to a cluster selected uniformly at **random** and then take the **centroids** of these initial clusters as the initial centers.

Bonus for CSCI 4372 and 5372 Students [10 points]; Mandatory for CSCI 6397 Students:

In addition, to “random partition”, implement the “maximin method”. This method chooses the first center **arbitrarily** from the data points and the remaining $(K - 1)$ centers are chosen successively as follows. In iteration i ($i = 2, 3, \dots, K$), the i -th center is chosen to be the data point with the greatest squared Euclidean distance to the nearest previously selected $(i - 1)$ centers. In other words, center 2 is chosen to be the farthest data point to center 1. Center 3 is chosen to be the data point with the greatest distance to the nearer of centers 1 and 2, that is, data point $\mathbf{x}$ with the greatest $\min(d(\mathbf{x}, \mathbf{c}_1), d(\mathbf{x}, \mathbf{c}_2))$ value, where $\mathbf{c}_1$ and $\mathbf{c}_2$ are the first and second centers, respectively, ‘d’ is the squared Euclidean distance, and the function “min(a, b)” returns the smaller of ‘a’ and ‘b’. Center i is chosen to be the data point $\mathbf{x}$ with the greatest $\min(d(\mathbf{x}, \mathbf{c}_1), d(\mathbf{x}, \mathbf{c}_2), \dots, d(\mathbf{x}, \mathbf{c}_{i-1}))$ value. By using additional memory, maximin can be implemented in $O(NDK)$ time, where N , D , and K denote the numbers of points, attributes, and clusters, respectively. A naïve implementation, however, requires $O(NDK^2)$ time. If you opt for a naïve implementation, you **may** have difficulty completing your runs.

Hint #3: ‘Arbitrary’ does **not** necessarily mean ‘random’. A random decision is arbitrary, but an arbitrary decision is **not** necessarily random.

Hint #4: Maximin works with both (ordinary) Euclidean and squared Euclidean distances. In both cases, the result is the same. Therefore, there is **no** good reason to use the computationally expensive ordinary Euclidean distance.

Bonus for CSCI 4372 and 5372 Students [10 points]; Mandatory for CSCI 6397 Students:

How does the random partition method compare against the random selection method in theory? Include your answer as a comment at the top of your main source file.

If your initialization method is randomized, execute it R times and tabulate the best result for each performance measurement separately (best initial SSEs, best final SSEs, and best # iterations) for each data set. Note that the best initial SSE, best final SSE, and best # iterations may be produced in different runs. In other words, the run that gives the best initial SSE will **not** necessarily be the run that gives the best final SSE. If your initialization method is deterministic (non-random), on the other hand, a simple table of the data sets and the corresponding measurements will suffice. There are several ways to tabulate this kind of data; try to find an intuitive way.

Output: Microsoft Excel or Apache OpenOffice Calc table(s) of the performance measurements for each combination of normalization and initialization methods. In other words, $\{10 \text{ data sets}\} \times \{1 \text{ or } 2 \text{ normalization methods}\} \times \{2 \text{ or } 3 \text{ initialization methods}\} \times \{3 \text{ performance measures}\} = 60 \text{ or } 180 \text{ values}$. Of course, you can split this table into smaller tables. For example, you can put 6 or 18 values into one table that represents a particular data set.

Language: C, C++, or Java. You may **only** use the built-in facilities of these languages. In other words, you may **not** use third-party libraries.

Submission: Submit your source code (C, C++, or Java source files) and output file(s) through Blackboard. Do **not** submit your entire project folder or the input data files I provided to you (I already have them 😊). In addition, do **not** submit your files individually; pack them into a single archive (e.g., zip) and submit the archive file.

Grading: *Functional correctness* and *adherence to good programming practices* are respectively worth **90%** and **10%** of your grade. Given a valid input, a functionally correct program is one that produces the correct output. There are a lot of generic or language-specific guides on good programming practices on the WWW. You **must** identify an appropriate one, include its URL at the top of your program source code(s), and use it throughout this project. In general, you should pay more attention to structural issues (e.g., modularity) than cosmetic ones (e.g., naming and formatting). Remember that if your program is incorrect, whether or not you adhered to good programming practices is immaterial.

Warning #2: There is a discussion forum for this phase on Blackboard. As indicated in the syllabus, **frequent** and **meaningful** participation in discussion forums is part of your overall grade. You are strongly encouraged to subscribe to this forum so that you receive email notifications of posts. Remember **not** to post code fragments (or URLs pointing to code fragments) on the forum. In other words, the discussions should be about ideas and algorithms, **not** about their implementation.

Hint #5: Common problems in this phase mainly stem from the careless reading of this document. Read the document **at least** twice before working on your program. Then, read it **at least** once after you are done.